

# Cache Invalidation using Rabbitmq

## Introduction:

Caching is a fundamental technique for improving system performance by storing frequently accessed data closer to the application. However, in distributed systems where multiple instances of an application are running, maintaining cache consistency becomes challenging. Cache invalidation is the process of removing or updating cached data to ensure data integrity and synchronization across distributed systems. This documentation explores cache invalidation strategies and demonstrates the implementation of a cache invalidation system using RabbitMQ, a popular message broker.

### Understanding Cache Invalidation

Cache invalidation is crucial in distributed systems to maintain data consistency across multiple servers. When the underlying data changes, cache entries need to be invalidated to prevent serving stale or incorrect data. Cache invalidation involves identifying and selectively invalidating cache entries based on certain criteria or events. Common cache invalidation strategies include time-based expiration, key-based invalidation, and event-based invalidation.

## Cache Invalidation Challenges in Distributed Systems:

Cache invalidation in distributed systems introduces several challenges that need to be addressed:

- **Consistency and Synchronization:** Ensuring cache consistency across multiple servers can be complex. When one server invalidates a cache entry, the invalidation message must propagate reliably and efficiently to other servers to ensure consistent cache state.
- **Scalability and Performance:** Cache invalidation mechanisms must be designed to handle large-scale distributed systems while maintaining optimal performance. As the system grows, the cache invalidation process should scale horizontally without introducing performance bottlenecks.

## Introducing RabbitMQ for Cache Invalidation:

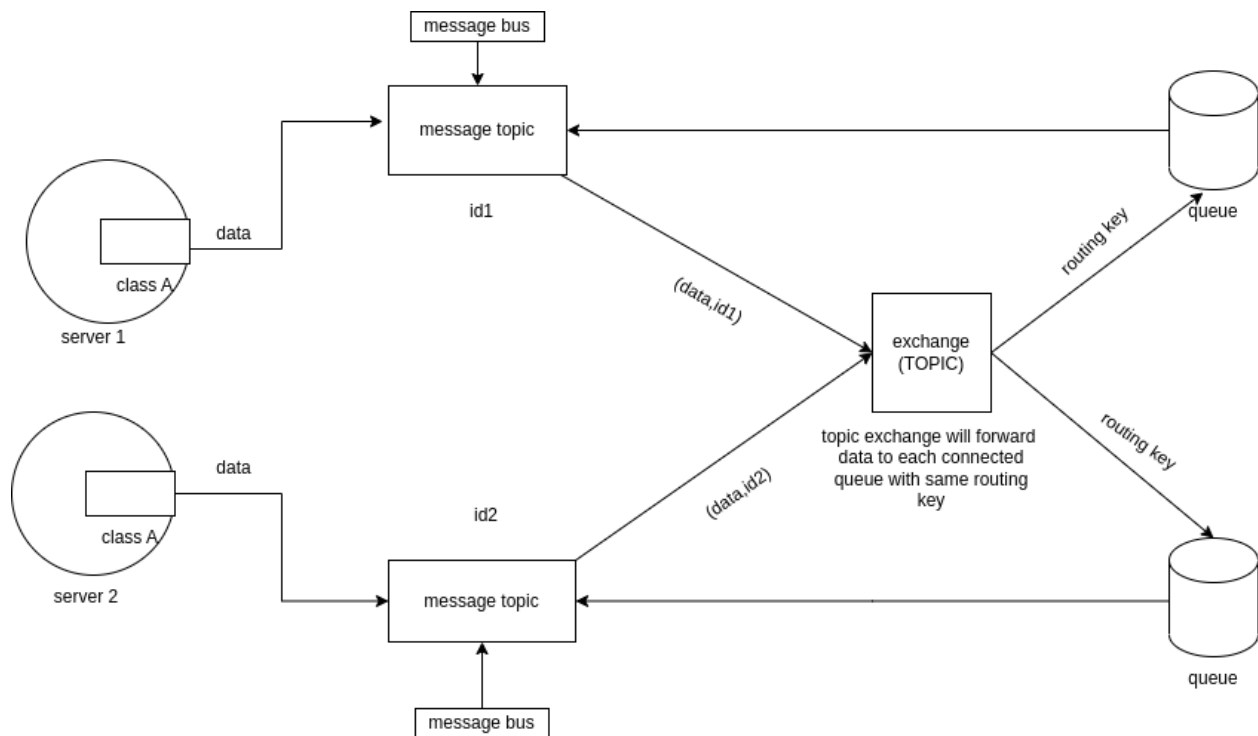
RabbitMQ is a robust and widely used message broker that can be leveraged for building cache invalidation systems in distributed environments. It provides the following benefits:

- **Decoupling:** RabbitMQ enables asynchronous communication between cache invalidators, allowing them to operate independently and at their own pace. Cache invalidation messages can be published to RabbitMQ without requiring immediate processing, decoupling the cache invalidation process from the application logic.
- **Reliability:** RabbitMQ ensures reliable message delivery through message acknowledgments and persistent storage. Messages can be durably stored in RabbitMQ even if the consumers are temporarily offline, ensuring that cache invalidation messages are not lost.
- **Scalability:** RabbitMQ supports distributed deployments, enabling cache invalidators to scale horizontally as the system grows. It can handle high message throughput and

efficiently distribute messages across multiple nodes.

## Designing:

1. **Create Message Bus:** The first step is to create a message bus, which acts as a central component for cache invalidation. The message bus is responsible for managing the connection with the message broker (such as RabbitMQ) and generating a unique ID for each class or component that needs cache invalidation.
2. **Create Message Topic:** Once the message bus is set up, you can create a cache topic. A cache topic represents a specific topic or channel for cache invalidation messages. When creating a Message topic, you need to provide the topic name and the data type associated with cache invalidation.
3. **Publish Cache Invalidation Data:** With the Message topic created, you can publish cache invalidation data to the associated queue. This can be done by invoking the publish method of the cache topic and providing the cache key or relevant data that needs to be invalidated.
4. **Consume Cache Invalidation Data:** The cache topic also allows consuming cache invalidation data from the associated queue. This is done by registering a callback function. The callback function is executed asynchronously when a cache invalidation message is received, allowing you to handle the cache invalidation logic.
5. **Cache Invalidation Function:** Within the callback function, you will invoke the cache invalidation function or method to invalidate the cache based on the received data. This is the place where you can implement your specific cache invalidation logic, such as removing or updating cache entries.
6. **Unique ID Checking:** To prevent duplicate cache invalidation actions, it is essential to incorporate a unique ID check. Each class or component that needs cache invalidation will have a unique ID generated by the message bus. When receiving cache invalidation data, you can compare the ID associated with the received data with the ID of the current class or component. If they match, you can ignore the data as it has already been invalidated.



## Example:

Let's assume we have two servers, s1 and s2, connected to RabbitMQ through the message bus. Both servers have the same class, ClassA, and the same routing key.

### 1. ClassA in s1 Server Invalidation:

- ClassA in s1 performs a cache invalidation operation for a specific key, let's say x.
- ClassA publishes the key x to the cache topic associated with the routing key.
- The cache topic, using the message bus, sends the cache invalidation data with the routing key to the topic exchange.

### 2. Publishing to Topic Exchange:

- The topic exchange receives the cache invalidation data with the routing key.
- The topic exchange publishes the data to all queues that are bound to the same routing key.

### 3. Consuming Cache Invalidation Data:

- Each server, s1 and s2, has its individual queue bound to the routing key.
- The cache topic in each server consumes the cache invalidation data from their respective queues.

### 4. Checking Unique ID and Invalidating Cache:

- Upon receiving the cache invalidation data, the cache topic in each server checks the unique ID associated with the received data.

- If the unique ID matches the ID of the message topic (e.g., id1 for s1 and id2 for s2), it means the cache invalidation has already been processed by the current server.
- In this case, the cache topic in the server discards the data from the queue, as it has already been invalidated.
- If the unique ID doesn't match the ID of the message topic, it means the cache invalidation is intended for the current server, and further cache invalidation logic can be performed.

## **Conclusion:**

Cache invalidation is a critical aspect of maintaining cache consistency in distributed systems. By leveraging RabbitMQ as a message broker, developers can implement an efficient and scalable cache invalidation system. This documentation provided an overview of cache invalidation strategies, discussed the challenges in distributed systems, and demonstrated the implementation of a cache invalidation system using RabbitMQ.